

Flight Delay Prediction

Nivedhitha D

Abstract. Flights are said to be delayed when they arrive later than the scheduled arrival time. This delay is predominantly influenced by environmental conditions. Flight delay is vexatious for passengers and also incurs an agonizingly high financial loss to airlines and countries. A structured prediction system is an indispensable tool that can help aviation authorities effectively alleviate flight delays. This project aims to build a two-stage machine learning engine to effectively predict the arrival delay of a flight in minutes after departure based on real-time flight and weather data. A classifier first predicts if the flight will be delayed or not and subsequently a regression model predicts the arrival delay in minutes if the flight is expected to be delayed.

Keywords: Machine Learning · Two-stage Model · Flight Delay Prediction.

1 Introduction

Flight delay is extremely troublesome to passengers and aviation authorities alike. Apart from the disruption of the schedule of the involved parties, flight delays cause monumental financial losses to the airline company. To accommodate the unforeseen delay in the arrival of a flight, a reallocation of airport resources, impromptu crew management and a redraft of flight schedules may arise. In some cases, the airline may be required to compensate the passengers for the delay. To address this issue, this project aims to design a two-stage machine learning engine to predict the arrival delay of flights accurately. Flight delay prediction can be viewed as a pipe-lined operation of two sequential tasks: predicting whether a flight will be delayed or not (classification) and if the flight is delayed, to predict the arrival delay in minutes (regression). The model is based on a data set synthesized from all flights in the 15 selected airports in the USA from 2016 to 2017 and corresponding weather data wherein flights with an arrival delay greater than 15 minutes are categorized as delayed. The performance of various classification and regression models is studied and compared before constructing the pipe-lined engine.

2 Data Set

The flight data set is a comma-separated value file containing the details of the flight schedules and their on-time performance in all the airports in the USA for the years 2016 and 2017. The weather data is a set of json files containing weather data that was recorded every one hour over 2016 and 2017 for 15 airports in the USA. The airports for which the weather data is available are selected and the corresponding flight data is merged with the weather data based on the Origin, Destination, date and time attributes. The time attribute of the flight data is rounded off to the nearest hour before merging with the weather data. The raw data set is then subject to data cleaning to handle missing data and redundant attributes. The processed data set consists of 18,51,436 data points. The models designate 80 per cent of the data points for training and the remaining 20 per cent of the data points for testing.

Table 1. The airports for which weather data is available.

| | | | | |
|-----|-----|-----|-----|-----|
| ATL | CLT | DEN | DFW | EWR |
| IAH | JFK | LAS | LAX | MCO |
| MIA | ORD | PHX | SEA | SFO |

Table 2. The weather data attributes considered.

| | | | |
|---------------|---------------|-------------|-----------|
| WindSpeedKmph | WindDirDegree | WeatherCode | precipMM |
| Visibility | Pressure | Cloudcover | DewPointF |
| WindGustKmph | tempF | WindChillF | Humidity |
| date | time | airport | |

Table 3. The flight schedule and performance attributes considered.

| | | | |
|-----------------|-----------------|-----------------|------------|
| FlightDate | Quarter | Year | Month |
| DayofMonth | DepTime | DepDel15 | CRSDepTime |
| DepDelayMinutes | OriginAirportID | DestAirportID | ArrTime |
| CRSArrTime | ArrDel15 | ArrDelayMinutes | |

3 Classification

3.1 Overview

Classification is the first stage of the pipe-lined model and aims to predict whether a scheduled flight will be delayed or not. Flights that are delayed have the target variable 'ArrDel15' set to 1 and those which are on time have the target variable 'ArrDel15' set to 0. The performance of five different models was studied and compared, namely, Logistic Regression, Decision Trees, Random Forest, Extra Trees and Gradient Boosting.

3.2 Performance Metrics

| | Actually Positive (1) | Actually Negative (0) |
|---------------------------|--|---|
| Predicted Positive (1) | True Positives (TPs) Flights correctly predicted to be delayed | False Positives (FPs) Flights wrongly predicted to be delayed |
| Predicted Negative (0) | False Negatives (FNs) Flights wrongly predicted to be on time | True Negative (TNs) Flights correctly predicted to be on time |

Fig. 1. Confusion Matrix.

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

$$Precision = \frac{TP}{TP + FP} \quad (2)$$

$$Recall = \frac{TP}{TP + FN} \quad (3)$$


$$F1 - score = \frac{2 * Precision * Recall}{Precision + Recall} \quad (4)$$

Table 4. Area Under Receiver Operating Characteristic Curve.

| AUC Value | Inference |
|-------------------|--|
| AUC = 1 | The classifier is able to perfectly distinguish between all the positive and the negative class points correctly |
| AUC = 0 | The classifier is predicting all negatives as positives, and all positives as negatives |
| AUC = 0.5 | The classifier is unable to distinguish the positive and negative class points thereby predicting a random or constant class for all the data points |
| Between 0.5 and 1 | There is a high chance that the classifier will be able to distinguish i.e. the classes more numbers of TPs and TNs than FNs and FPs |

Table 5. Results from the different classification models.

| Algorithm | Precision | | Recall | | F1-Score | | Accuracy |
|---------------------|-----------|------|--------|------|----------|------|----------|
| | 0 | 1 | 0 | 1 | 0 | 1 | |
| Logistic Regression | 0.92 | 0.89 | 0.98 | 0.68 | 0.95 | 0.77 | 0.92 |
| Decision Tree | 0.92 | 0.68 | 0.91 | 0.71 | 0.92 | 0.69 | 0.87 |
| Extra Trees | 0.93 | 0.82 | 0.96 | 0.73 | 0.94 | 0.77 | 0.91 |
| Gradient Boosting | 0.92 | 0.90 | 0.88 | 0.68 | 0.95 | 0.78 | 0.92 |
| Random Forest | 0.93 | 0.86 | 0.97 | 0.70 | 0.95 | 0.78 | 0.92 |

**Fig. 2.** Area under ROC for the different classification models.

3.3 Class Imbalance

It can be observed that the models were better at predicting the recall and F1 scores of the negative class (class 0) when compared to the scores of the positive class (class 1). The poor performance of the classifiers on class 1 relative to class 0 on the data set is because of the inherent skew towards the class 'Not-Delayed' flights. This skew can lead to incorrect learning and misleading optimistic performance (accuracy) because the information is biased to one class. Out of 18,51,436 data points, only 3,88,058 data points are delayed flights.




Fig. 3. Imbalanced data set class distribution.

3.4 Overcoming Imbalance

To overcome this bias, we need to perform sampling. There are two sampling methods to make the data set balanced:

- **Under-sampling**

The majority class is under-sampled to ensure balance by randomly selecting examples from the majority class to delete from the training data set repeatedly until the desired class distribution is achieved.

- **Over-sampling**

The minority class is over-sampled to ensure balance by synthesizing data points from the minority class from the training data set repeatedly until the desired class distribution is achieved.

To preserve the existing data, over-sampling was employed to balance the data set. **Synthetic Minority Over-sampling TEchnique (SMOTE)** was used to sample the data set as it synthesises data points that have smooth variation and high correlation with the existing data set.




Fig. 4. Class distribution after SMOTE.

3.5 Classifier Performance Comparison after SMOTE

Table 6. Results from the different classification models after SMOTE.

| Algorithm | Precision | | Recall | | F1-Score | | Accuracy |
|---------------------|-----------|------|--------|------|----------|------|----------|
| | 0 | 1 | 0 | 1 | 0 | 1 | |
| Logistic Regression | 0.94 | 0.74 | 0.93 | 0.78 | 0.93 | 0.76 | 0.90 |
| Decision Tree | 0.92 | 0.66 | 0.90 | 0.71 | 0.91 | 0.68 | 0.86 |
| Extra Trees | 0.94 | 0.77 | 0.94 | 0.76 | 0.94 | 0.76 | 0.86 |
| Gradient Boosting | 0.93 | 0.80 | 0.95 | 0.75 | 0.94 | 0.77 | 0.91 |
| Random Forest | 0.93 | 0.81 | 0.95 | 0.74 | 0.94 | 0.78 | 0.91 |

SMOTE oversampling improves the performance of the models by increasing the recall of class 1. Since the F1 Score is the weighted harmonic mean of precision and recall, it was used as the primary criterion to choose the Random Forest Classifier, having the highest F1 Score (0.78).




Fig. 5. Area under ROC for the different classification models after SMOTE.

4 Regression

4.1 Overview

Regression is the second stage of the pipe-lined model and aims to predict the arrival delay in minutes if the flight is classified as ‘Delayed’ by the classifier. The flights having ‘ArrDelayMinutes’ greater than 15 were used to train the regression model. The performance of five different models was studied and compared, namely, Logistic Regression, Decision Trees, Random Forest, Extra Trees and Gradient Boosting.

4.2 Performance Metrics

$$MeanSquaredError(MSE) = \frac{1}{n} \sum_{i=1}^n (y_i - x_i)^2 \quad (5)$$

$$RootMeanSquaredError(RMSE) = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - x_i)^2} \quad (6)$$

$$MeanAbsoluteError(MAE) = \frac{1}{n} \sum_{i=1}^n |y_i - x_i| \quad (7)$$

$$R^2 Score = 1 - \frac{\sum_{i=1}^n (y_i - x_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (8)$$

R-squared Score is a measure of the ability of a model to predict the variances in the data set accurately. The Random Forest Regressor having the most promising R-squared Score (0.937) and RMSE (15.038) was chosen.

Table 7. Results from the different regression models.

| Algorithm | MSE | RMSE | MAE | R^2 Score |
|-------------------|---------------|-----------|-----------|-------------|
| Linear Regression | 242.641989 | 15.576970 | 10.590204 | 0.933416 |
| Decision Tree | 473.420932 | 21.758238 | 14.584725 | 0.870087 |
| Extra Trees | 229.312907 | 15.143081 | 10.465352 | 0.937073 |
| Gradient Boosting | 230.90.909551 | 15.195708 | 10.323951 | 0.936635 |
| Random Forest | 226.167177 | 15.038856 | 10.384938 | 0.937937 |

4.3 Range-wise Analysis

In this section, the data set is split into ranges of arrival delay minutes and the performance of the pipe-lined Random Forest Regressor is studied in each range. The flight arrival delay ranges from 0 to 1210 minutes. The frequency distribution plot of the arrival delay indicates that the majority of data points are observed in the 0 - 200 range. The bar plot reveals that most of the data points have 'ArrDelayMinutes' ranging between 0 - 100 minutes. This explains the better performance of the model as we approach the maximum. As the range increases, the number of data points decreases, indicating that flights with very high flight delays are less. As a result of decreasing data points, the values of MSE, RMSE and MAE scores increase. The 1000-1210 range error is relatively small compared to the magnitude of the bin.




Fig. 6. Range-wise frequency analysis.

Table 8. Range-wise regression analysis.

| Range | MSE | RMSE | MAE | R^2 Score |
|-----------|------------|-----------|-----------|-------------|
| 0-100 | 96.720685 | 9.834668 | 6.634876 | 0.856493 |
| 100-200 | 237.147213 | 15.399585 | 9.194358 | 0.658591 |
| 200-500 | 339.491540 | 18.425296 | 10.801726 | 0.914323 |
| 500-1000 | 411.110167 | 20.275852 | 11.399412 | 0.982030 |
| 1000-1210 | 27.836775 | 5.276057 | 4.457500 | 0.995092 |

5 Pipe-lined Model

The flow chart represents the two-stage flight delay prediction machine learning model. The pipe-lined model involves chaining the best performing classifier before the best regressor. The data was reprocessed and trained to perform classification using the Random Forest Classifier. The Random Forest Classifier was chosen as it has the maximum F1 Score (0.78) and area under ROC (0.85). The flight delay needs to be calculated only for the flights that will be delayed. Thus, only those data points that were predicted to be delayed by the classifier are selected to perform regression and predict the flight arrival delay in minutes. The Random Forest Regressor was chosen as it has the highest R-squared score (0.937) and lower values of RMSE (15.038) and MAE (10.384).

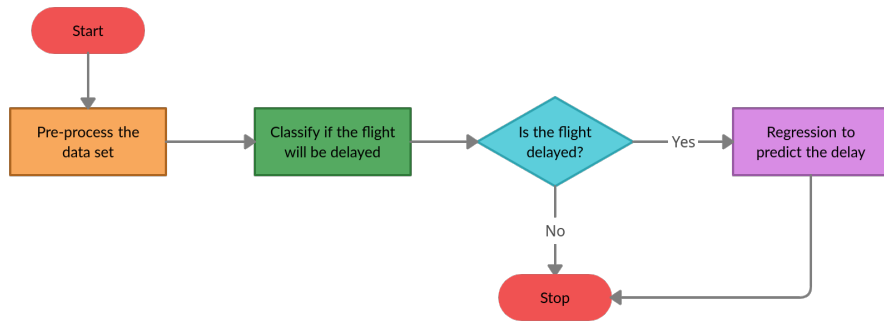


Fig. 7. Flight delay prediction as a pipe-lined operation of two sequential tasks: predicting whether a flight will be delayed or not (classification) and if the flight is delayed, to predict the arrival delay in minutes (regression).

Table 9. Pipe-lined model performance evaluation.

| Metric | Value |
|---------------|--------------|
| MSE | 127.311361 |
| RMSE | 11.283234 |
| MAE | 7.178468 |
| R^2 Score | 0.977392 |

6 Result

The flight and weather data were combined into a single data set and pre-processed to train a two-stage machine learning model that predicts flight arrival delay. Due to class imbalance, there was an inherent bias towards the majority class, 'Not Delayed' flights (class 0). The data was sampled using SMOTE before classification to overcome the bias. Out of five different algorithms, the Random Forest classifier gave the best F1 score (0.78) and Recall (0.74) for the delayed flights. Subsequently, the Random Forest regressor was pipe-lined, giving MAE 7.178 minutes and RMSE 11.283 minutes with an R2 score of 0.977. In conclusion, the flight delay prediction was efficient and the Machine Learning model exhibited good performance.